



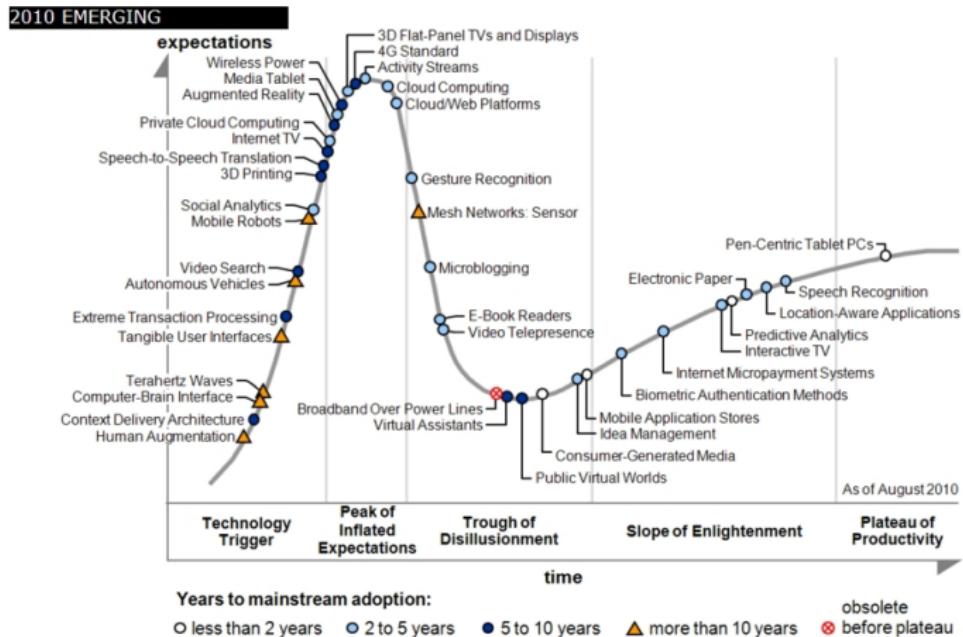
# Introducing the Cloud

Gwendal Simon

<http://perso.telecom-bretagne.eu/gwendalsimon>  
@gwendal

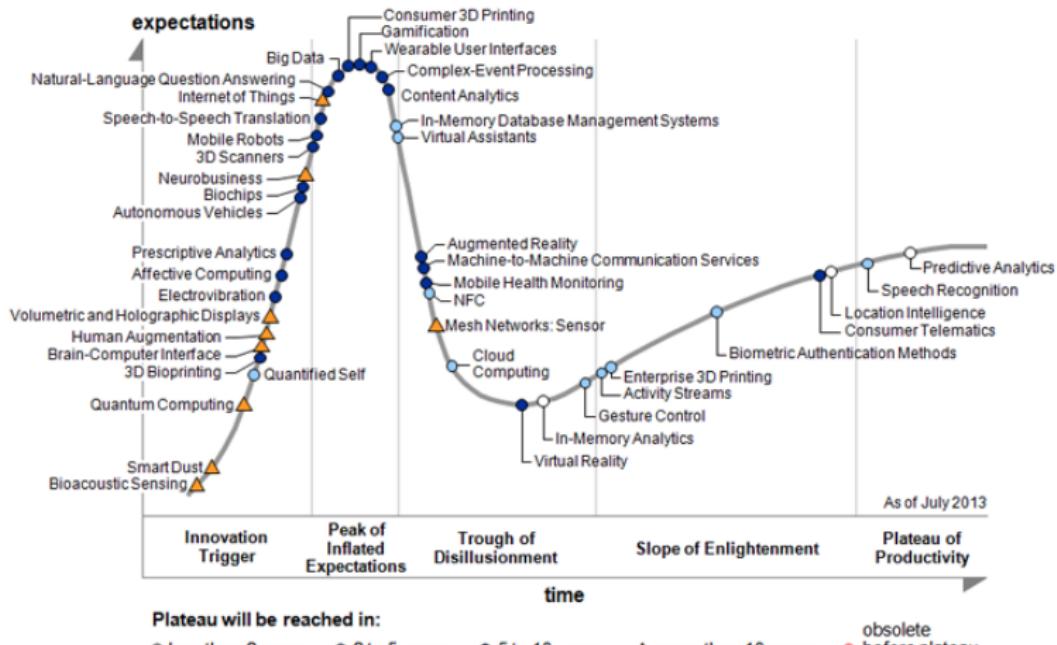


# Already hype when you entered TB



Gartner

# Hype in a different way now





# Avancement

- 1 Mini Introduction
- 2 Case Studies
- 3 Basic Elements of a Datacenter
- 4 Software Infrastructure
- 5 Conclusion



# In a nutshell



## In a nutshell

Cloud computing is a model for enabling convenient, *on-demand* network access to a *shared* pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly *provisioned* and released with *minimal management effort* or service provider interaction

NIST : <http://csrc.nist.gov/groups/SNS/cloud-computing/index.html>



# The \*aaS paradigm



# The \*aaS paradigm

## ■ SaaS : Software as a Service

- applications for *end-users*
  - email, office suite, photos sharing, video storage



# The \*aaS paradigm

- SaaS : Software as a Service
  - applications for *end-users*
    - email, office suite, photos sharing, video storage
- PaaS : Platform as a Service
  - services for *web app developers*
    - workflow facilities and various basic services



# The \*aaS paradigm

- SaaS : Software as a Service
  - applications for *end-users*
    - email, office suite, photos sharing, video storage
- PaaS : Platform as a Service
  - services for *web app developers*
    - workflow facilities and various basic services
- IaaS : Infrastructure as a Service
  - resources for *developers*
    - servers, network equipment, memory, CPU



# Selected Pictures



# Selected Pictures





# Selected Pictures



# Selected Pictures





# Avancement

- 1 Mini Introduction
- 2 Case Studies
- 3 Basic Elements of a Datacenter
- 4 Software Infrastructure
- 5 Conclusion



# Case Studies

- A Commercial Offer : Amazon
- Switching to the Cloud : Netflix



# Amazon Web Services

- DynamoDB
- Simple Storage Service (S3)
- Elastic Compute Cloud (EC2)
- Elastic Load Balancing (ELB)
- Virtual Private Cloud (VPC)
- Glacier
- Elastic MapReduce (EMR)
- CloudFront



# Case Studies

- A Commercial Offer : Amazon
- Switching to the Cloud : Netflix



# See Slides

<http://www.slideshare.net/adrianco/netflix-global-cloud>



# Avancement

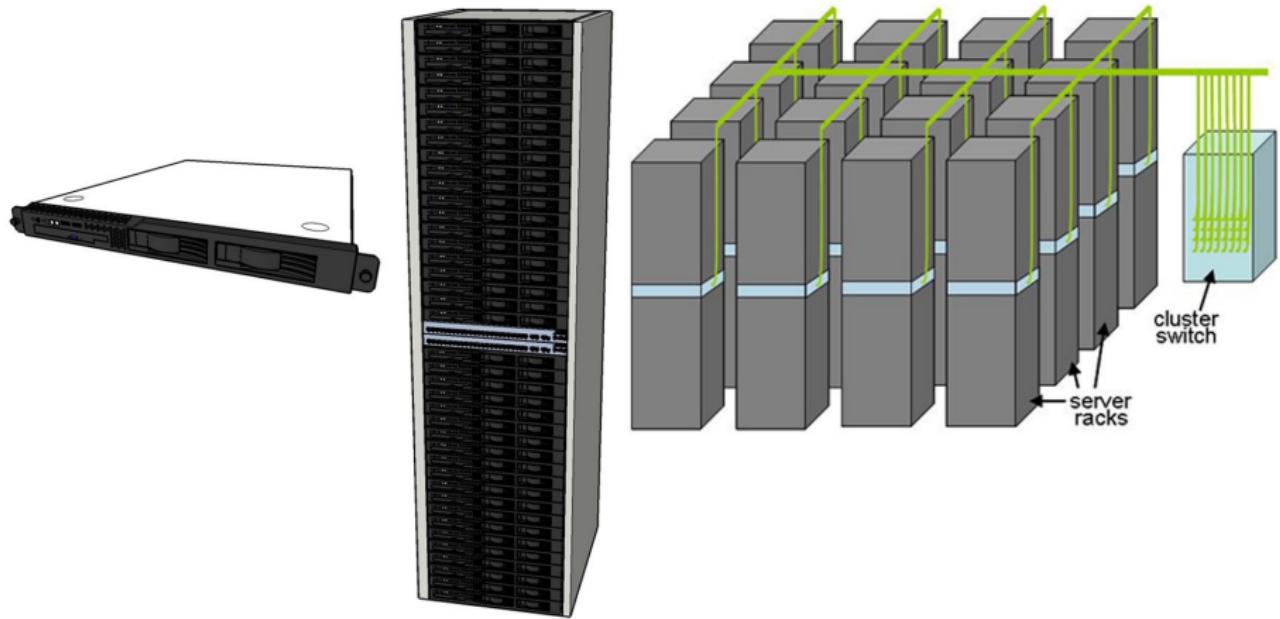
- 1 Mini Introduction
- 2 Case Studies
- 3 Basic Elements of a Datacenter
- 4 Software Infrastructure
- 5 Conclusion



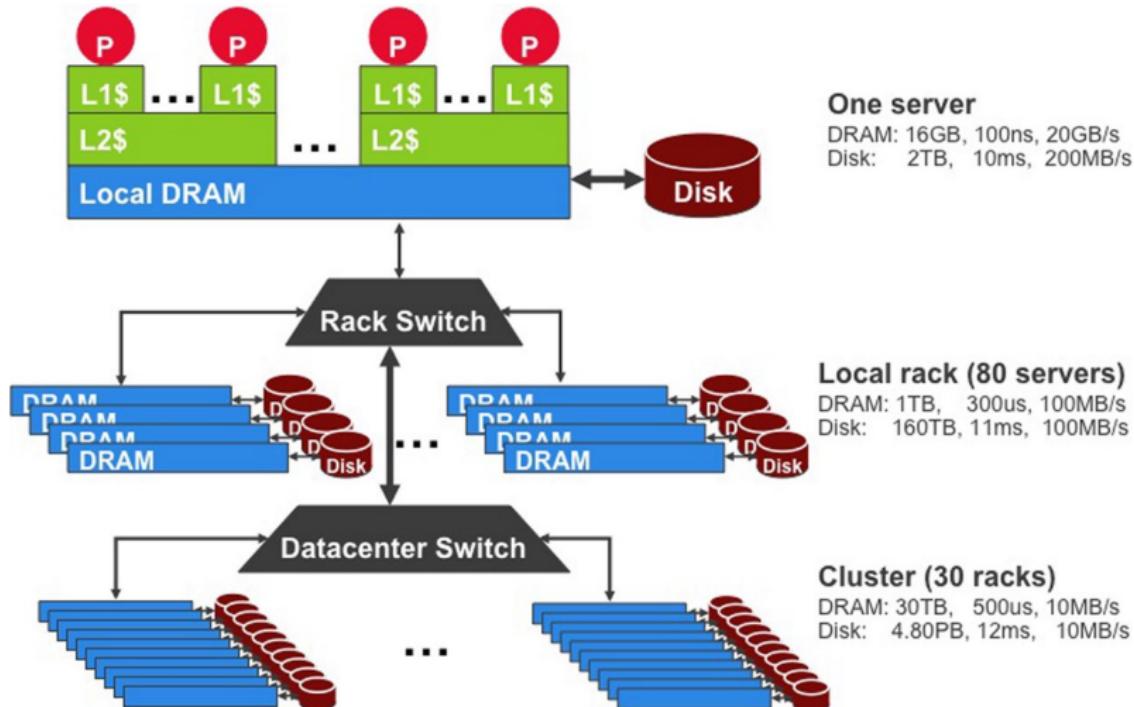
# Basic Elements of a Datacenter

- Computing Architecture
- Energy
- Dealing with Failures

# Architecture Basics



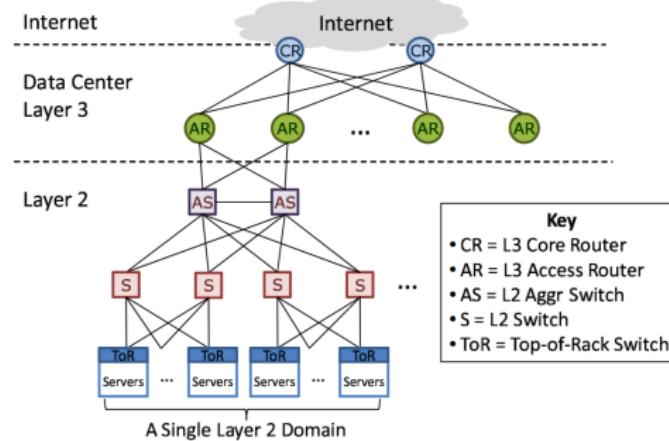
# Storage and Access Time (in 2009)



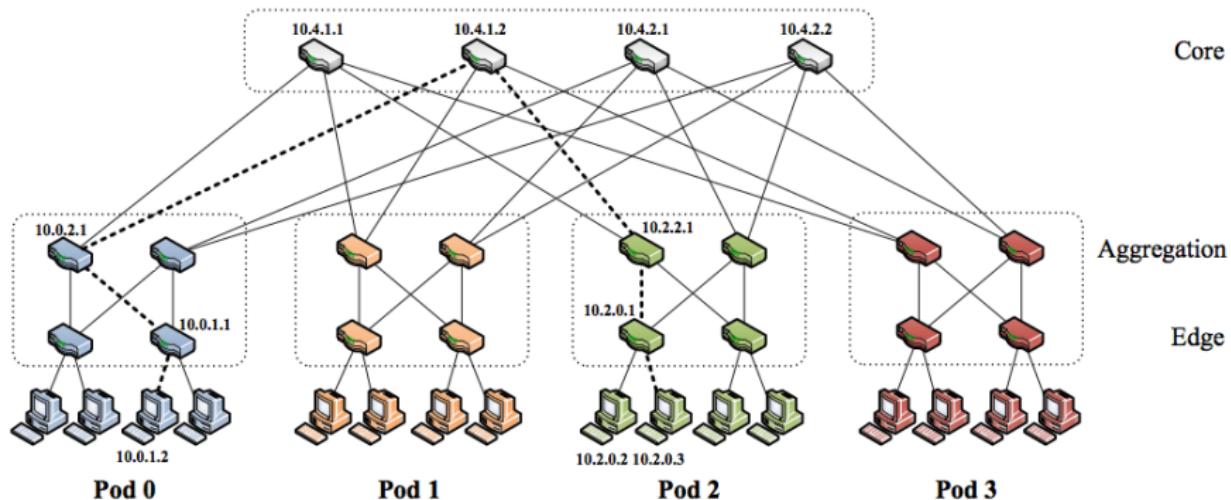


## Oversubscription :

worst-case communication between servers  
capacity of network topology



# Network Fat-Tree Topologies





# Typical Technologies and Prices (in 2011)

## Server :

- computation : 24-32 cores at \$200 per core
- network : \$150 per 10 Gbps port NIC



# Typical Technologies and Prices (in 2011)

## **Server :**

- computation : 24-32 cores at \$200 per core
- network : \$150 per 10 Gbps port NIC

## **10 Gbps switch :**

- today : 48-64 ports at \$450 per port
- trend : up to 150 ports at \$100 per port

# Typical Technologies and Prices (in 2011)

## Server :

- computation : 24-32 cores at \$200 per core
- network : \$150 per 10 Gbps port NIC

## 10 Gbps switch :

- today : 48-64 ports at \$450 per port
- trend : up to 150 ports at \$100 per port

## Memory :

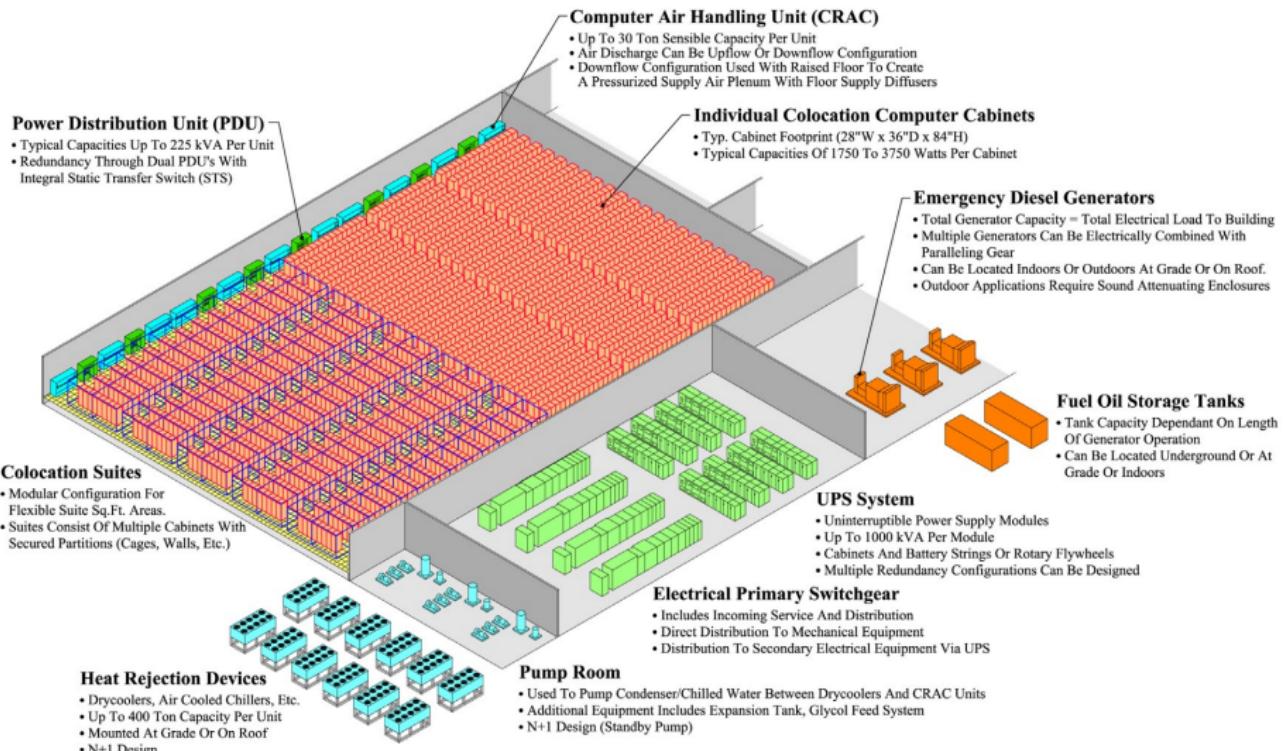
Technology	Access time (ns)	Cost (\$/MB)	Max. size
SRAM	4	27	≈ 250 Mbits
RDRAM	15	0.27	≈ 2 Gbits
DRAM	55	0.016	≈ 100 Gbits



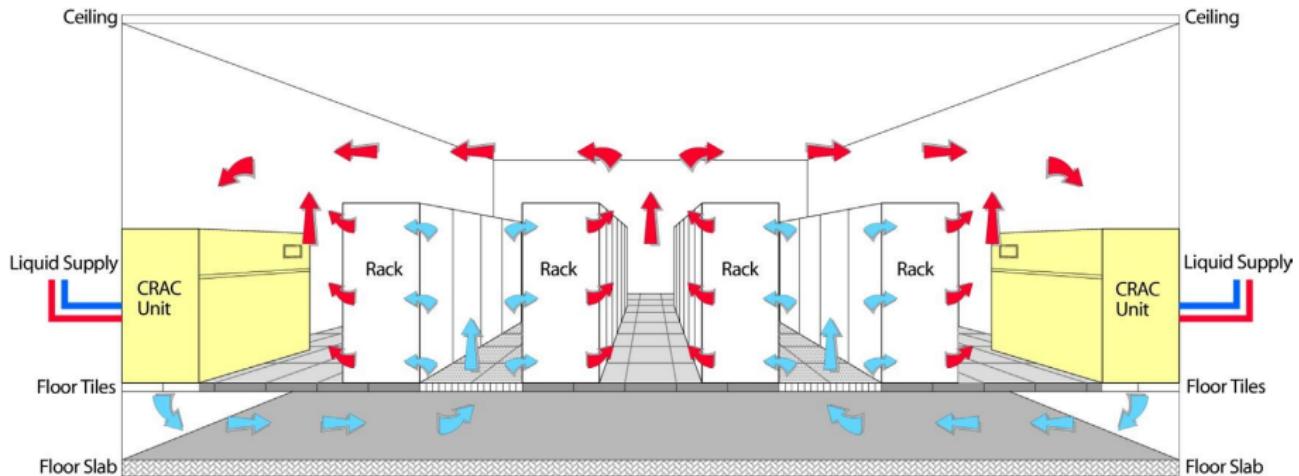
# Basic Elements of a Datacenter

- Computing Architecture
- Energy
- Dealing with Failures

# Main Electric Components



# Cooling Challenge





# Evaluating Energy Efficiency

*Power Usage Effectiveness* =  $\frac{\text{total energy}}{\text{energy in equipment}}$

- first generation datacenter PUE  $\geq 3.0$
- toward PUE around 1.1



# Evaluating Energy Efficiency

*Power Usage Effectiveness* =  $\frac{\text{total energy}}{\text{energy in equipment}}$

- first generation datacenter PUE  $\geq 3.0$
- toward PUE around 1.1

*Server PUE* =  $\frac{\text{total server power}}{\text{critical component power}}$

- basic SPUE is 1.7
- state-of-the-art servers reach 1.1



# Evaluating Energy Efficiency

*Power Usage Effectiveness* =  $\frac{\text{total energy}}{\text{energy in equipment}}$

- first generation datacenter PUE  $\geq 3.0$
- toward PUE around 1.1

*Server PUE* =  $\frac{\text{total server power}}{\text{critical component power}}$

- basic SPUE is 1.7
- state-of-the-art servers reach 1.1

and **computing efficiency** ?

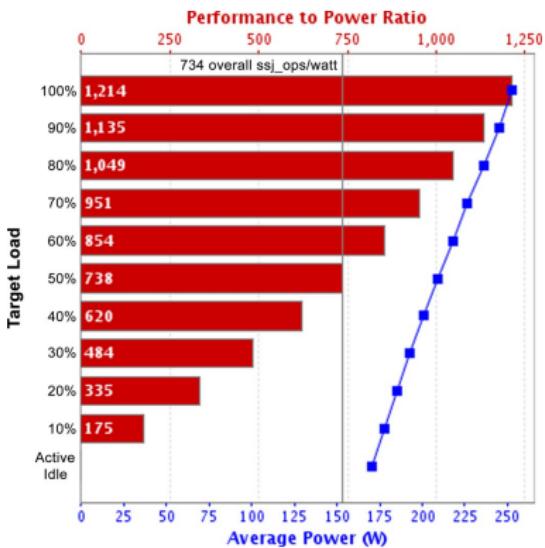


# Computing Efficiency

*Benchmarking cluster-level efficiency : on-going work*

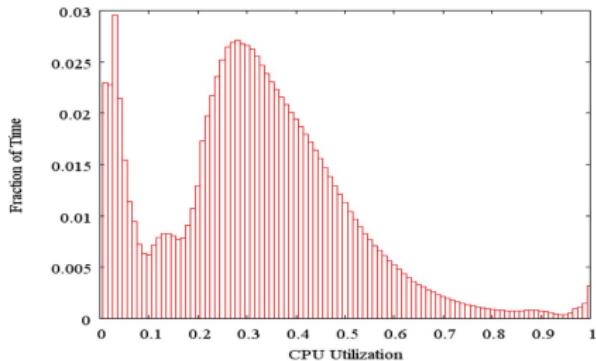
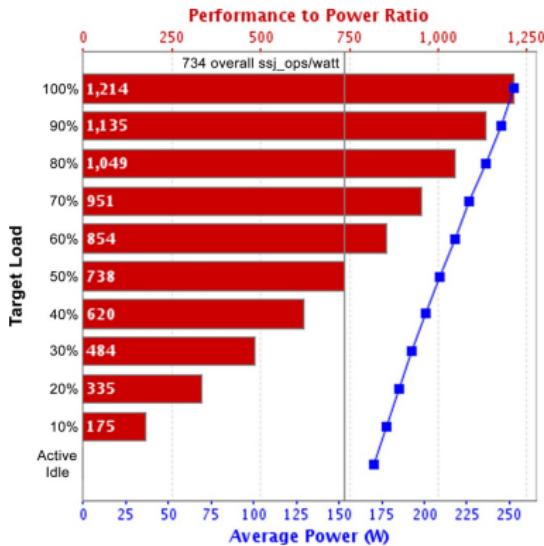
# Computing Efficiency

*Benchmarking cluster-level efficiency : on-going work*  
*Benchmarking individual computer is easier*

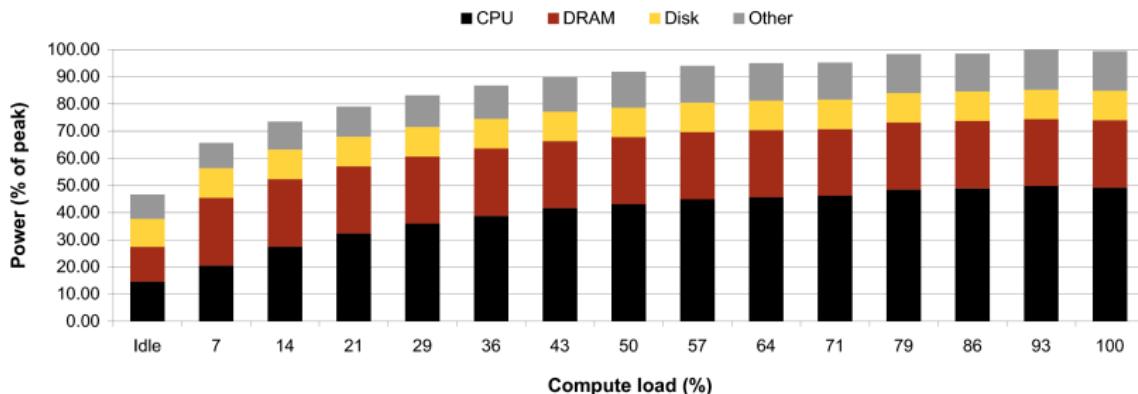


# Computing Efficiency

*Benchmarking cluster-level efficiency : on-going work*  
*Benchmarking individual computer is easier*



# Toward Energy-Proportional Servers

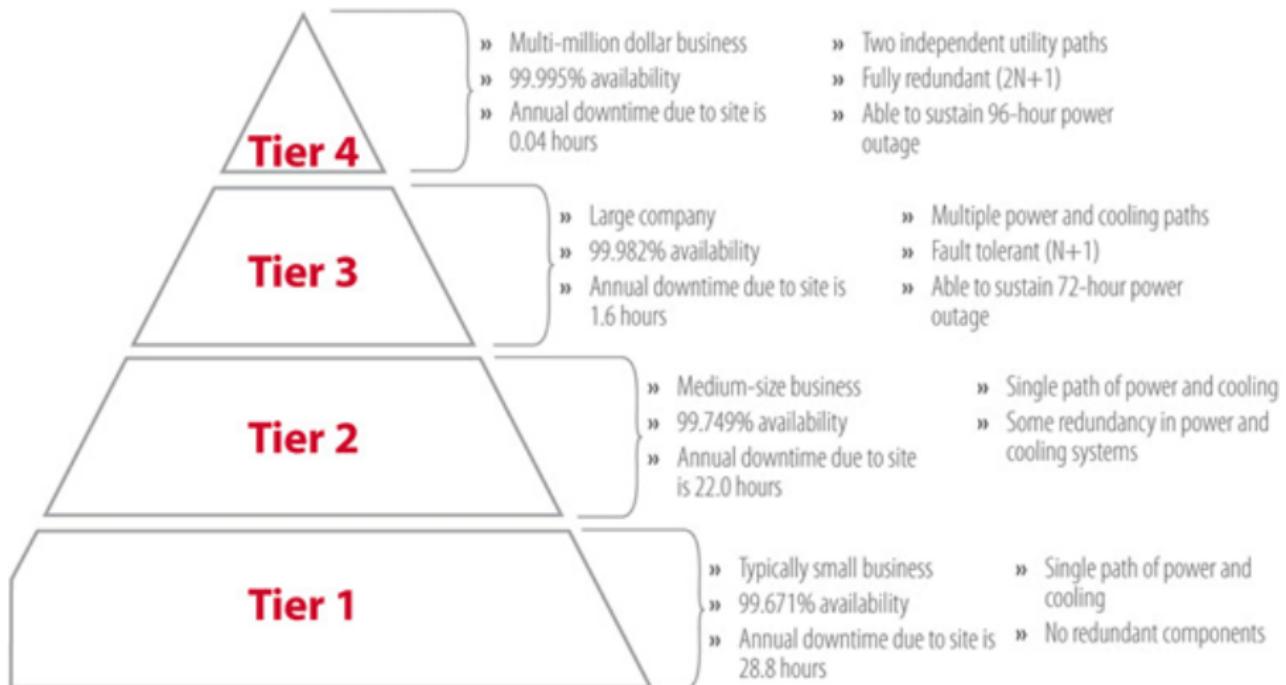




# Basic Elements of a Datacenter

- Computing Architecture
- Energy
- Dealing with Failures

# Data-Center Tiers





# Origin of Impacting Failures

Cause	% of events
software	33 %
configuration	28 %
human	13 %
network	12 %
hardware	11 %
other	3 %



# Origin of Impacting Failures

Cause	% of events
software	33 %
configuration	28 %
human	13 %
network	12 %
hardware	11 %
other	3 %

hardware faults are masked by fault-tolerant software



## Failures and Crash

Average machine availability is  $\simeq 99.9\%$

- 95% of machines restart less than once a month
- 80% of restart events last less than 10 minutes



## Failures and Crash

Average machine availability is  $\simeq 99.9\%$

- 95% of machines restart less than once a month
- 80% of restart events last less than 10 minutes

Software most frequent faults (in one year) :

- *DRAM soft-errors* : 1% experience uncorrectable errors
- *disk soft-errors* : 3% of drives see corrupted sectors



# Avancement

- 1 Mini Introduction
- 2 Case Studies
- 3 Basic Elements of a Datacenter
- 4 Software Infrastructure
- 5 Conclusion



# Software Infrastructure

- Cluster-Level : MapReduce



# Motivation

Map/Reduce is a software framework for easily writing applications which **process** vast amounts of data (multi-terabyte data-sets) *in-parallel* on large clusters (thousands of nodes) of commodity hardware in a reliable, fault-tolerant manner.



# Functional Programming

Two fundamentals functions :

- **map** : apply a function to a list of elements

- $\text{map } f \ [ ] = [ ]$

- $| \text{map } f \ [x :: xs] = (f \ x) :: (\text{map } f \ xs)$

- $\text{map square } [1,2,5] \rightarrow [1,4,25]$



# Functional Programming

Two fundamentals functions :

- **map** : apply a function to a list of elements

- $\text{map } f \ [ ] = [ ]$   
|  $\text{map } f \ [x :: xs] = (f \ x) :: (\text{map } f \ xs)$
- $\text{map square } [1,2,5] \rightarrow [1,4,25]$

- **reduce** : build a value from a function and a list

- $\text{reduce } f \ a \ [ ] = a$   
|  $\text{reduce } f \ a \ [x :: xs] = \text{reduce } f \ (f \ x \ a) \ xs$
- $\text{reduce add 0 } [1,3,6] \rightarrow 10$



# MapReduce

Implementing two functions w.r.t data (key, val)

■ **map** : smaller sub-problems distributed to nodes

- map (inKey, inVal)  $\rightarrow$  list (outKey, v)
- produces *intermediate* values with an output key



# MapReduce

Implementing two functions w.r.t data (key, val)

■ **map** : smaller sub-problems distributed to nodes

- map (inKey, inVal)  $\rightarrow$  list (outKey, v)
- produces *intermediate* values with an output key

■ **reduce** : combines results of sub-problems

- reduce (outKey, list v)  $\rightarrow$  outVal
- produces an output value from intermediate values



## Example : Word Count

```
map(filename, content):  
    for each w in content:  
        emitInt(w, 1)
```



# Example : Word Count

```
map(filename, content):  
    for each w in content:  
        emitInt(w, 1)  
  
reduce(word, partCount):  
    int result = 0  
    for pc in partCount:  
        result += pc  
    emit(result)
```

# Example : Word Count

```
map(filename, content):  
    for each w in content:  
        emitInt(w, 1)  
  
reduce(word, partCount):  
    int result = 0  
    for pc in partCount:  
        result += pc  
    emit(result)
```

```
map(file1, "hello me, goodbye me")→  
<hello,1><me,1><goodbye,1><me,1>  
  
map(file2, "hello you, bye you")→  
<hello,1><you,1><bye,1><you,1>
```



## Example : Word Count

```
map(filename, content):  
    for each w in content:  
        emitInt(w, 1)  
  
reduce(word, partCount):  
    int result = 0  
    for pc in partCount:  
        result += pc  
    emit(result)
```

map(file1, "hello me, goodbye me")→  
<hello,1><me,1><goodbye,1><me,1>

map(file2, "hello you, bye you")→  
<hello,1><you,1><bye,1><you,1>

a given *key* is allocated to a given *server*



## Example : Word Count

```
map(filename, content):  
    for each w in content:  
        emitInt(w, 1)  
  
reduce(word, partCount):  
    int result = 0  
    for pc in partCount:  
        result += pc  
    emit(result)
```

map(file1, "hello me, goodbye me") →  
<hello,1><me,1><goodbye,1><me,1>  
  
map(file2, "hello you, bye you") →  
<hello,1><you,1><bye,1><you,1>  
  
a given *key* is allocated to a given *server*  
  
reduce(hello,<1,1>) → 2  
...



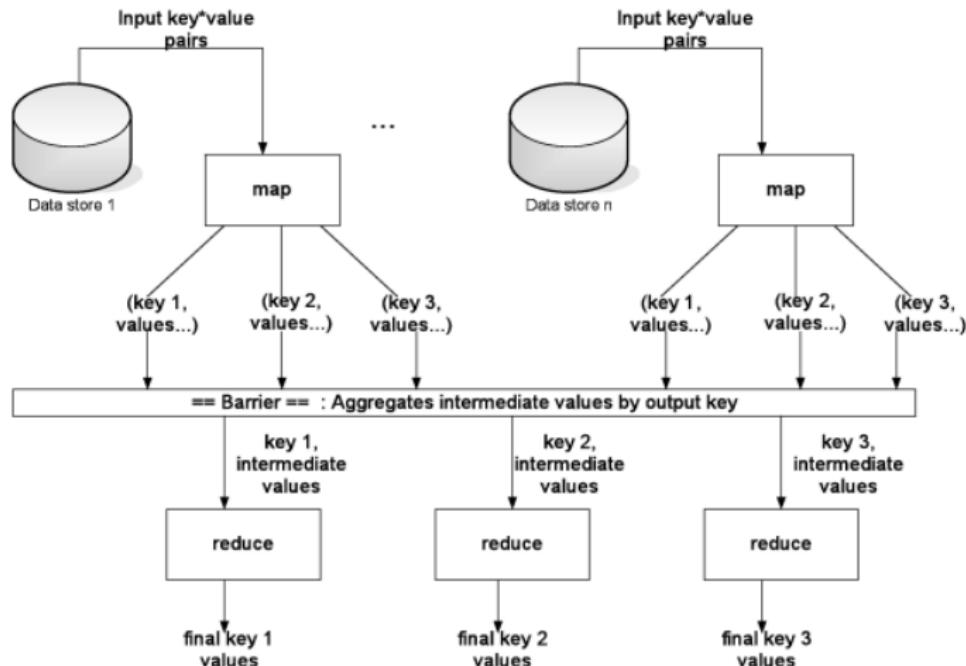
## Example : Word Count

```
map(filename, content):  
    for each w in content:  
        emitInt(w, 1)  
  
reduce(word, partCount):  
    int result = 0  
    for pc in partCount:  
        result += pc  
    emit(result)
```

map(file1, "hello me, goodbye me") →  
<hello,1><me,1><goodbye,1><me,1>  
  
map(file2, "hello you, bye you") →  
<hello,1><you,1><bye,1><you,1>  
  
a given *key* is allocated to a given *server*  
  
reduce(hello,<1,1>) → 2  
  
...

<hello,2><me,2><you,2><goodbye,1><bye,1>

# MapReduce Implemented





# Avancement

- 1 Mini Introduction
- 2 Case Studies
- 3 Basic Elements of a Datacenter
- 4 Software Infrastructure
- 5 Conclusion



## A 30-sec conclusion

A new industrial era :

- computer science does really matter
- be fluent in at (least one) cloud tech. or die