



# Cloud Computing, Open Stack and Data-centers

Gwendal Simon

<http://perso.telecom-bretagne.eu/gwendalsimon>  
@gwendal





## Recommended Book

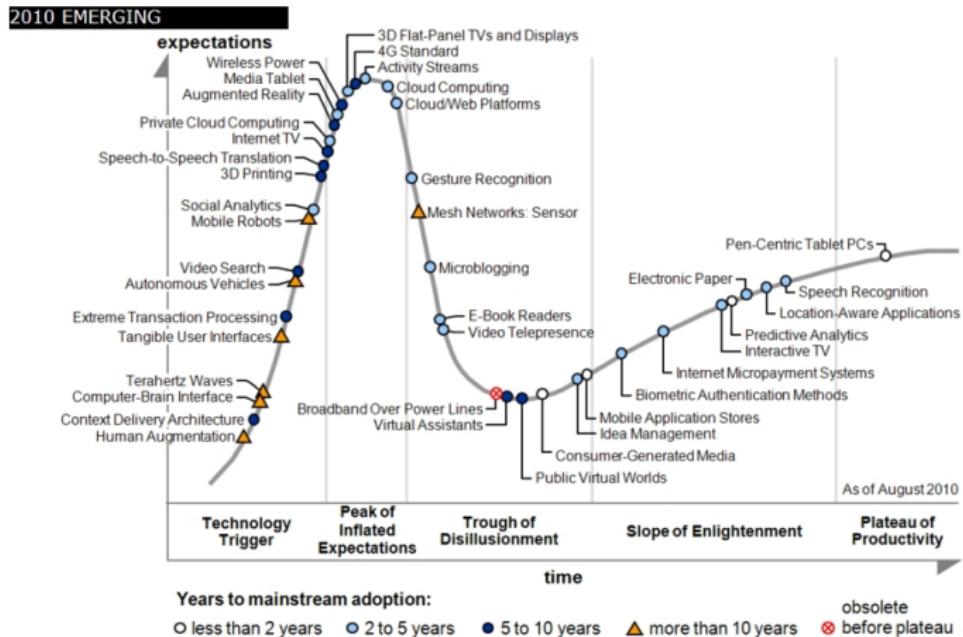
### **The Datacenter as a Computer : An Introduction to the Design of Warehouse-Scale Machines**

2nd edition (2013)

Luiz Barroso, Jimmy Clidaras, and Urs Hölzle

Morgan & Claypool publishers

# A hyped tech in 2010



Gartner

# No longer hyped, but operational

Figure 1. Hype Cycle for Emerging Technologies, 2015





## Definition in a nutshell

Cloud computing is a model for enabling convenient, *on-demand* network access to a *shared* pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly *provisioned* and released with *minimal management effort* or service provider interaction

NIST : <http://csrc.nist.gov/groups/SNS/cloud-computing/index.html>



# The \*aaS paradigm

- SaaS : Software as a Service
  - applications for *end-users*
    - email, office suite, photos sharing, video storage
- PaaS : Platform as a Service
  - services for *web app developers*
    - workflow facilities and various basic services
- IaaS : Infrastructure as a Service
  - resources for *developers*
    - servers, network equipment, memory, CPU



# Avancement

## 1 Basic Elements of a Datacenter

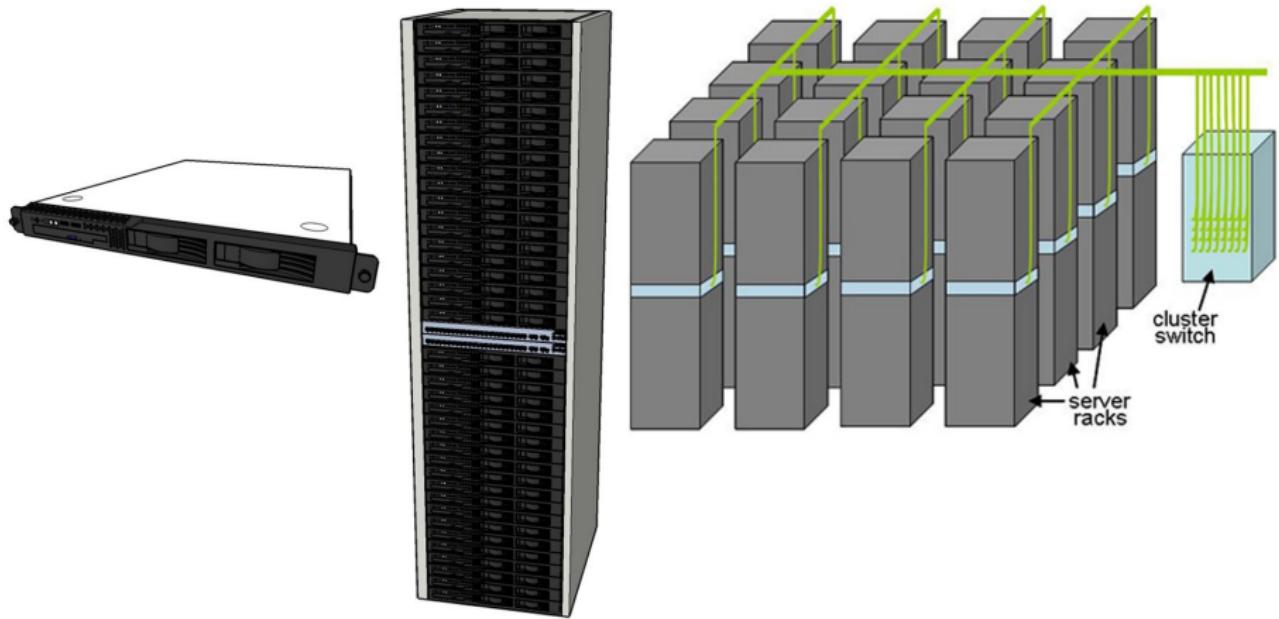
## 2 Case Studies

## 3 Dealing with Concerns

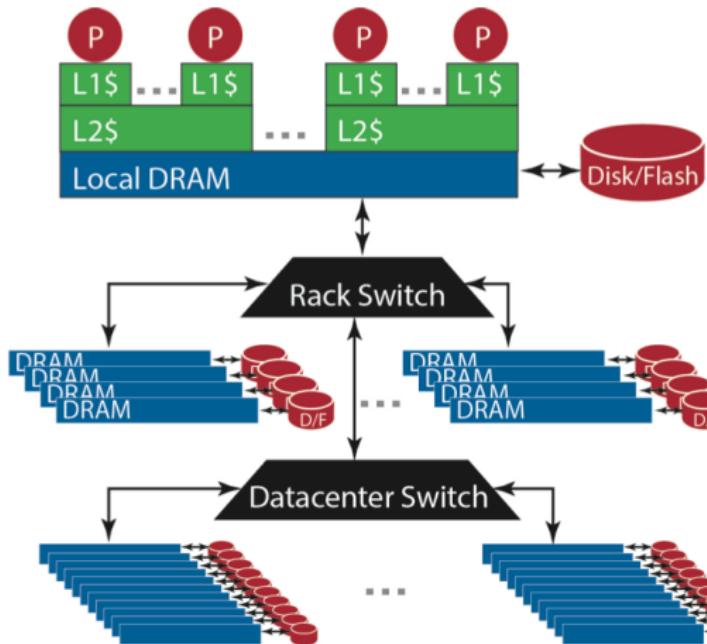
## 4 Software Infrastructure

## 5 Conclusion

# Architecture Basics



# Storage and Access Time



## One Server

DRAM: 16 GB, 100 ns, 20 GB/s  
Disk: 2T B, 10 ms, 200 MB/s  
Flash: 128 GB, 100 us, 1 GB/s

## Local Rack (80 servers)

DRAM: 1 TB, 300 us, 100 MB/s  
Disk: 160 TB, 11 ms, 100 MB/s  
Flash: 20 TB, 400 us, 100 MB/s

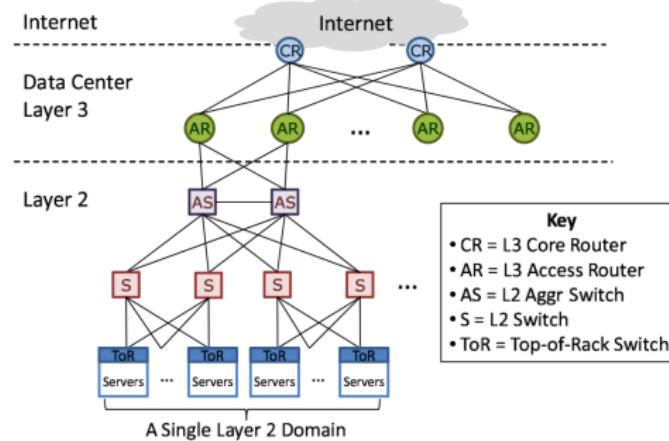
## Cluster (30 racks)

DRAM: 30 TB, 500 us, 10 MB/s  
Disk: 4.80 PB, 12 ms, 10 MB/s  
Flash: 600 TB, 600 us, 10 MB/s

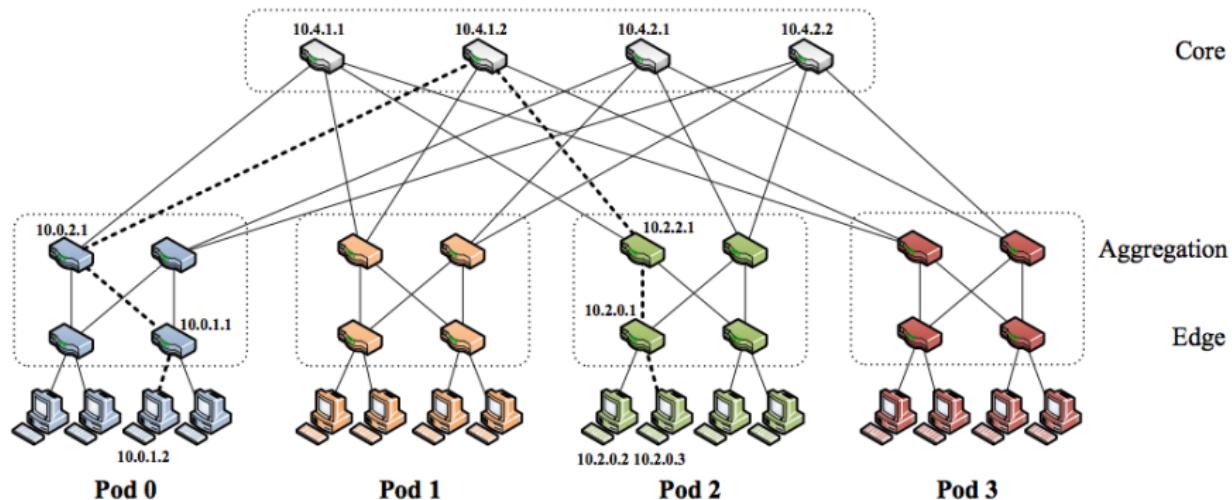


## Oversubscription :

worst-case communication between servers  
capacity of network topology



# Network Fat-Tree Topologies



# Typical Technologies and Prices (in 2013)

## Server :

- computation : 24-32 cores at \$200 per core
- network : \$150 per 10 Gbps port NIC

## 10 Gbps switch :

- today : 48-64 ports at \$450 per port
- trend : up to 150 ports at \$100 per port

## Memory :

Technology	Access time (ns)	Cost (\$/MB)	Max. size
SRAM	4	27	≈ 250 Mbits
RDRAM	15	0.27	≈ 2 Gbits
DRAM	55	0.016	≈ 100 Gbits



# Avancement

- 1 Basic Elements of a Datacenter
- 2 Case Studies
- 3 Dealing with Concerns
- 4 Software Infrastructure
- 5 Conclusion



# Case Studies

- A Commercial Offer : Amazon
- Switching to the Cloud : Netflix



# Amazon Web Services

- DynamoDB
- Simple Storage Service (S3)
- Elastic Compute Cloud (EC2)
- Elastic Load Balancing (ELB)
- Virtual Private Cloud (VPC)
- Glacier
- Elastic MapReduce (EMR)
- CloudFront



# Case Studies

- A Commercial Offer : Amazon
- Switching to the Cloud : Netflix



# See Slides

<http://www.slideshare.net/adrianco/netflix-global-cloud>



# Avancement

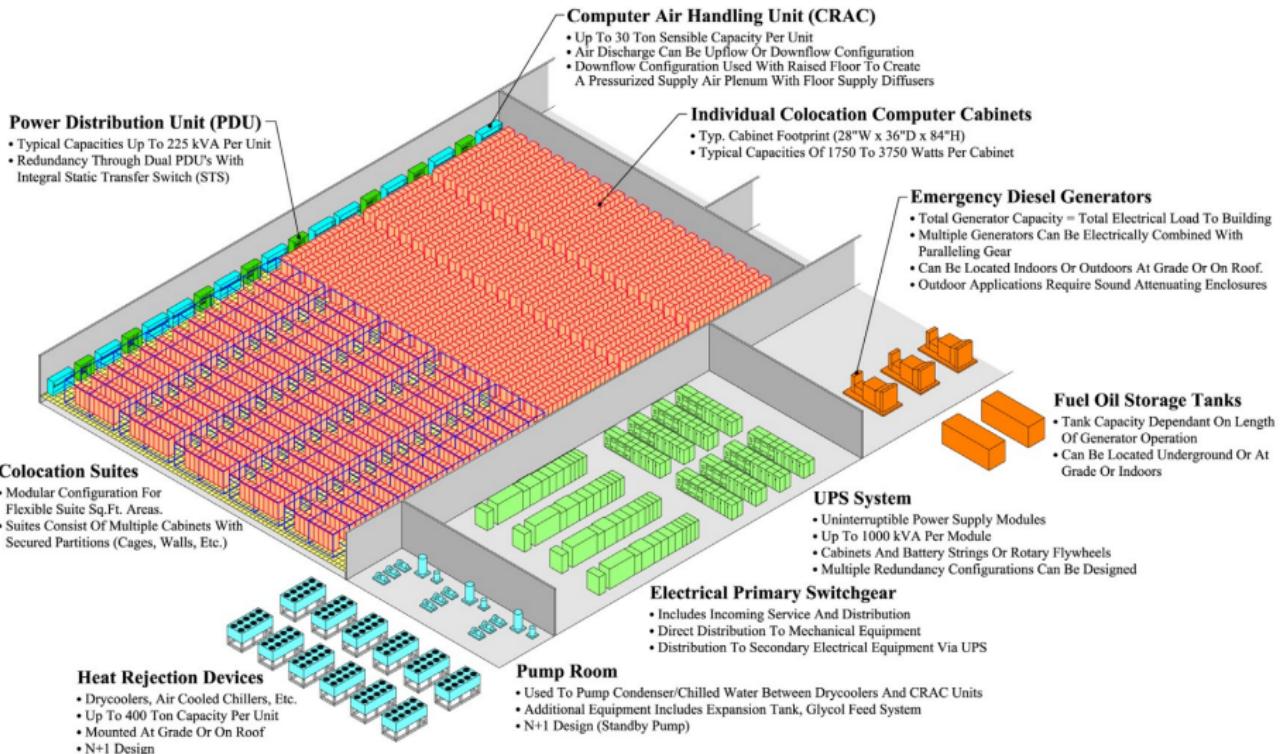
- 1 Basic Elements of a Datacenter
- 2 Case Studies
- 3 Dealing with Concerns
- 4 Software Infrastructure
- 5 Conclusion



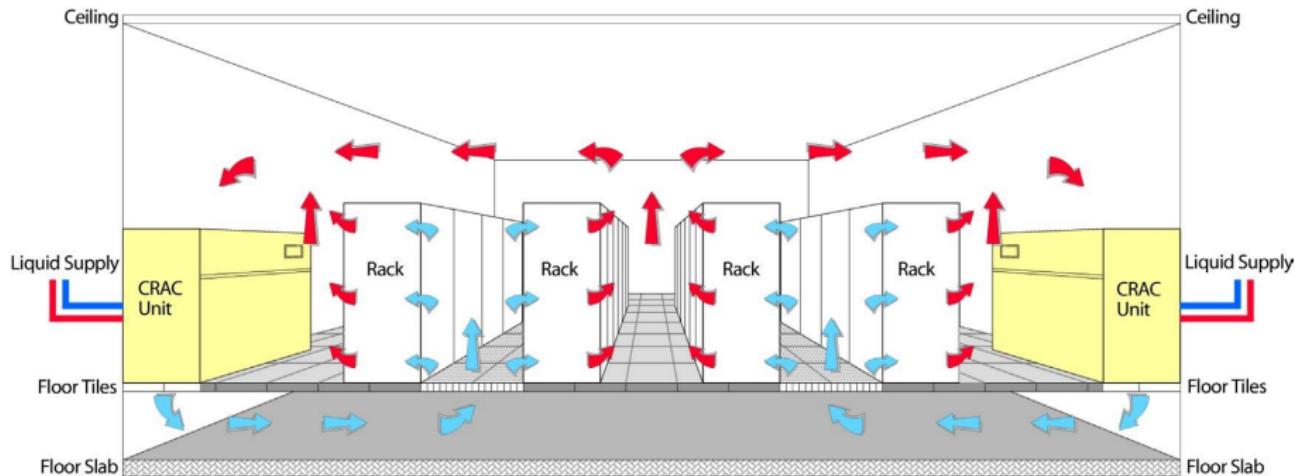
# Dealing with Concerns

- Energy
- Failures

# Main Electric Components

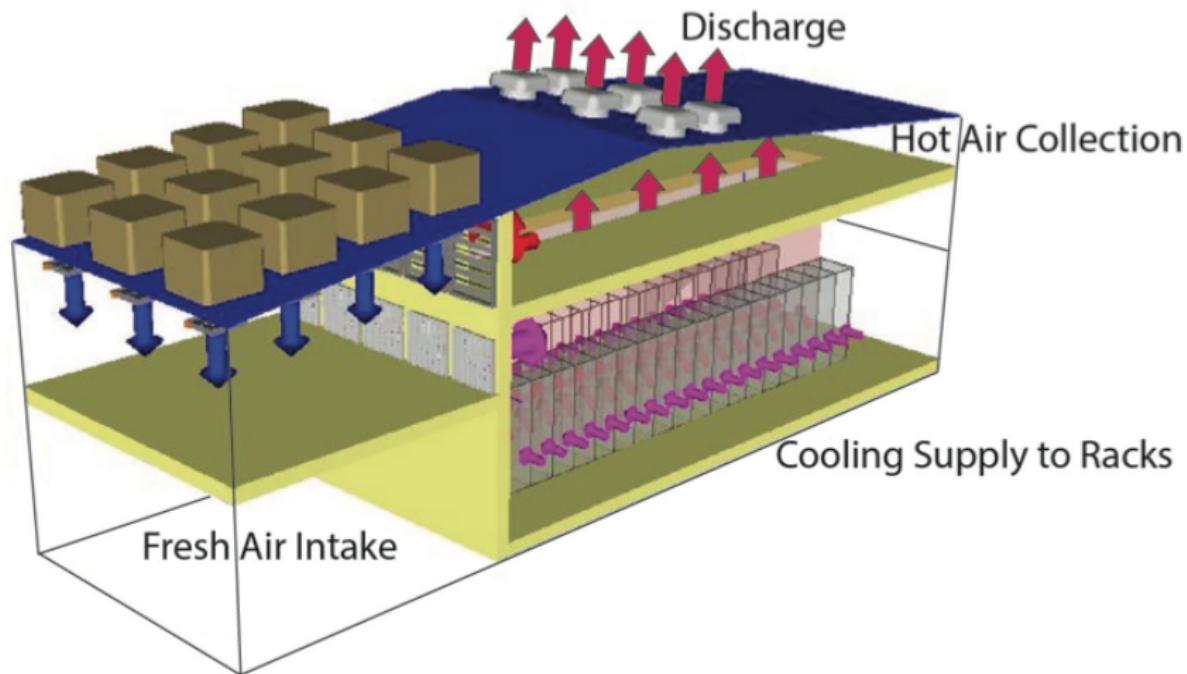


# Cooling Challenge (closed loop)





# Cooling Challenge (open loop)





# Evaluating Energy Efficiency

*Power Usage Effectiveness* =  $\frac{\text{total energy}}{\text{energy in equipment}}$

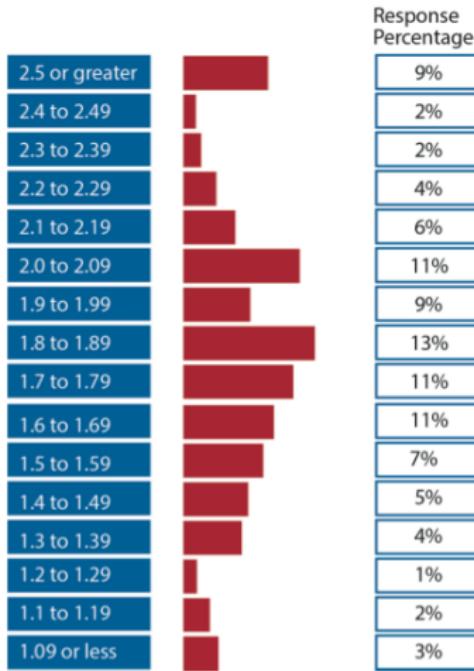
- first generation datacenter PUE  $\geq 3.0$
- toward PUE around 1.1

*Server PUE* =  $\frac{\text{total server power}}{\text{critical component power}}$

- basic SPUE is 1.7
- state-of-the-art servers reach 1.1

# Current data-center PUE

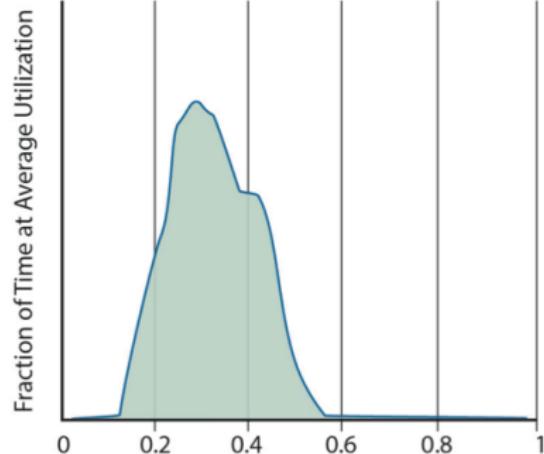
Average PUE of your largest data center:



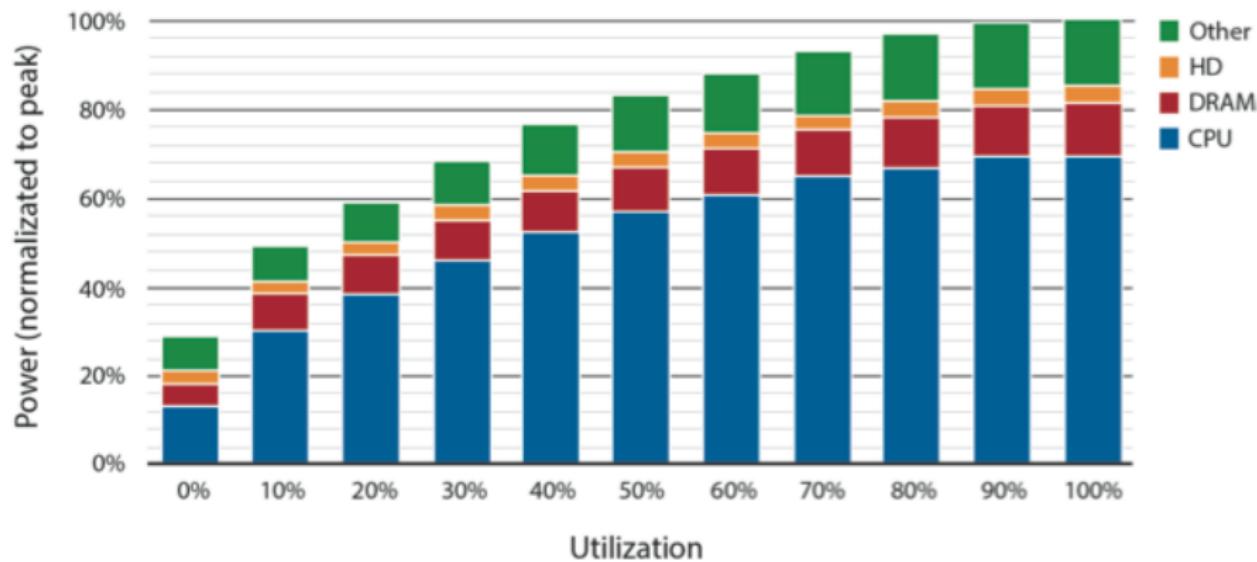
AVERAGE  
PUE  
1.8 – 1.89

# Computing Efficiency

*Benchmarking cluster-level efficiency : on-going work*  
*Benchmarking individual computer is easier*



# Toward Energy-Proportional Servers

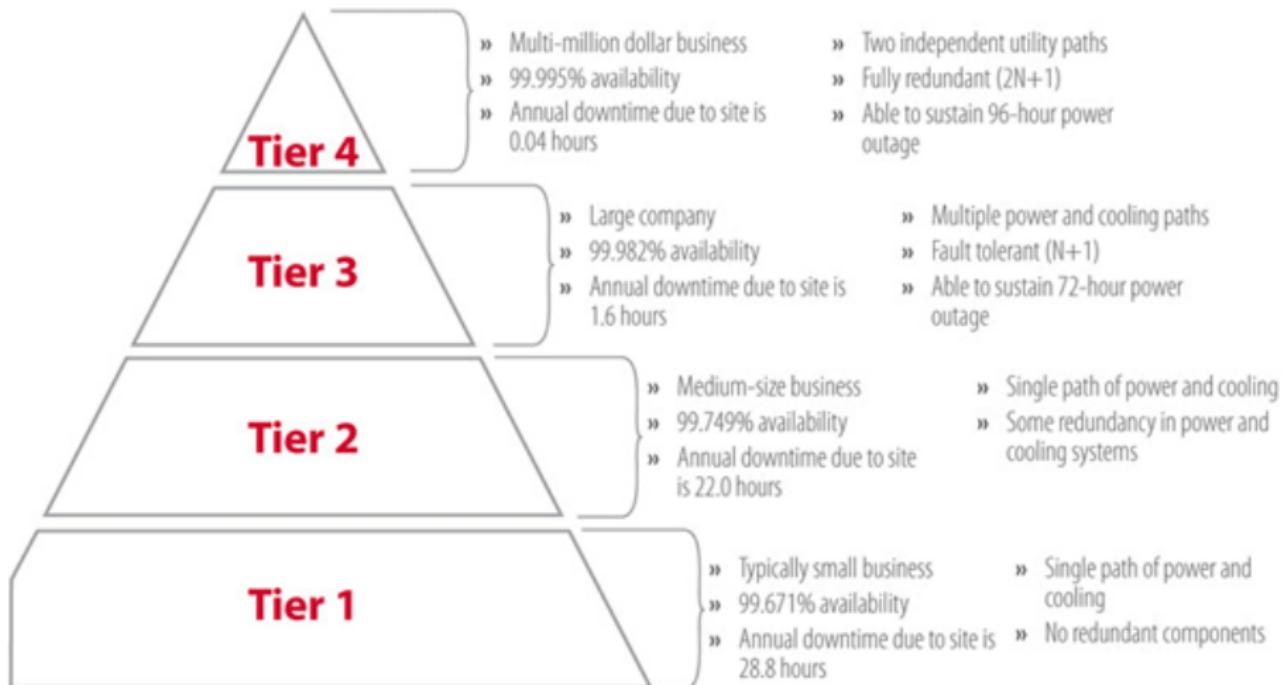




# Dealing with Concerns

- Energy
- Failures

# Data-Center Tiers





# Origin of Impacting Failures

Cause	% of events
software	33 %
configuration	28 %
human	13 %
network	12 %
hardware	11 %
other	3 %

hardware faults are masked by fault-tolerant software



## Failures and Crash

Average machine availability is  $\simeq 99.9\%$

- 95% of machines restart less than once a month
- 80% of restart events last less than 10 minutes

Software most frequent faults (in one year) :

- *DRAM soft-errors* : 1% experience uncorrectable errors
- *disk soft-errors* : 3% of drives see corrupted sectors



# Avancement

- 1 Basic Elements of a Datacenter
- 2 Case Studies
- 3 Dealing with Concerns
- 4 Software Infrastructure
- 5 Conclusion



# Software Infrastructure

- Cluster-Level : MapReduce
- Open Stack



# Motivation

Map/Reduce is a software framework for easily writing applications which **process** vast amounts of data (multi-terabyte data-sets) *in-parallel* on large clusters (thousands of nodes) of commodity hardware in a reliable, fault-tolerant manner.



# Functional Programming

Two fundamentals functions :

- **map** : apply a function to a list of elements

- $\text{map } f \ [ ] = [ ]$   
|  $\text{map } f \ [x :: xs] = (f \ x) :: (\text{map } f \ xs)$
- $\text{map square } [1,2,5] \rightarrow [1,4,25]$

- **reduce** : build a value from a function and a list

- $\text{reduce } f \ a \ [ ] = a$   
|  $\text{reduce } f \ a \ [x :: xs] = \text{reduce } f \ (f \ x \ a) \ xs$
- $\text{reduce add 0 } [1,3,6] \rightarrow 10$



# MapReduce

Implementing two functions w.r.t data (key, val)

■ **map** : smaller sub-problems distributed to nodes

- map (inKey, inVal)  $\rightarrow$  list (outKey, v)
- produces *intermediate* values with an output key

■ **reduce** : combines results of sub-problems

- reduce (outKey, list v)  $\rightarrow$  outVal
- produces an output value from intermediate values

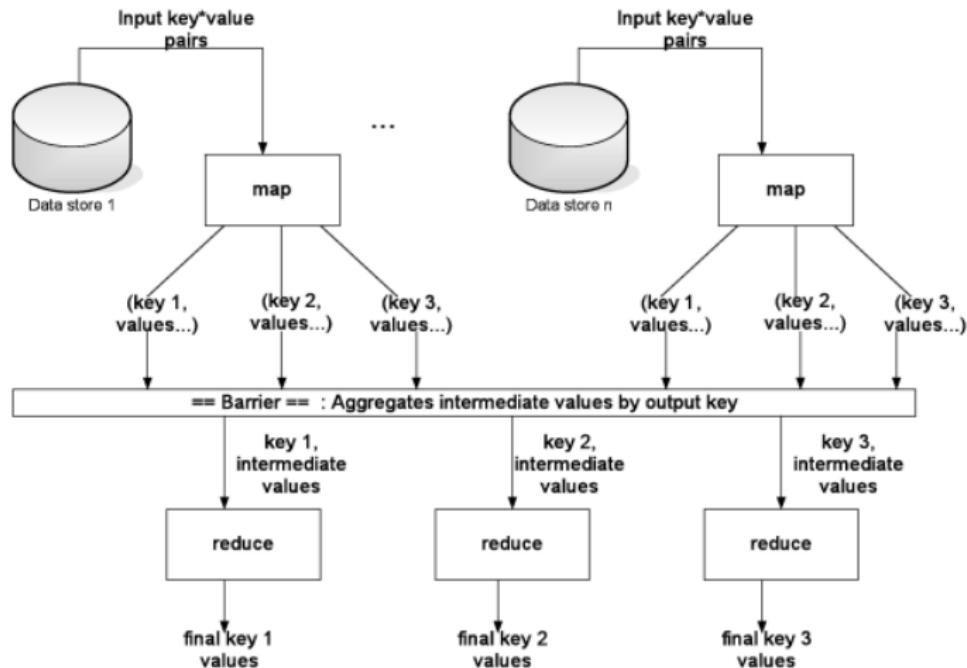
# Example : Word Count

```
map(filename, content):  
    for each w in content:  
        emitInt(w, 1)  
  
reduce(word, partCount):  
    int result = 0  
    for pc in partCount:  
        result += pc  
    emit(result)
```

map(file1, "hello me, goodbye me") →  
<hello,1><me,1><goodbye,1><me,1>  
  
map(file2, "hello you, bye you") →  
<hello,1><you,1><bye,1><you,1>  
  
a given *key* is allocated to a given *server*  
  
reduce(hello,<1,1>) → 2  
  
...

<hello,2><me,2><you,2><goodbye,1><bye,1>

# MapReduce Implemented





# Software Infrastructure

- Cluster-Level : MapReduce
- Open Stack



## Main Idea

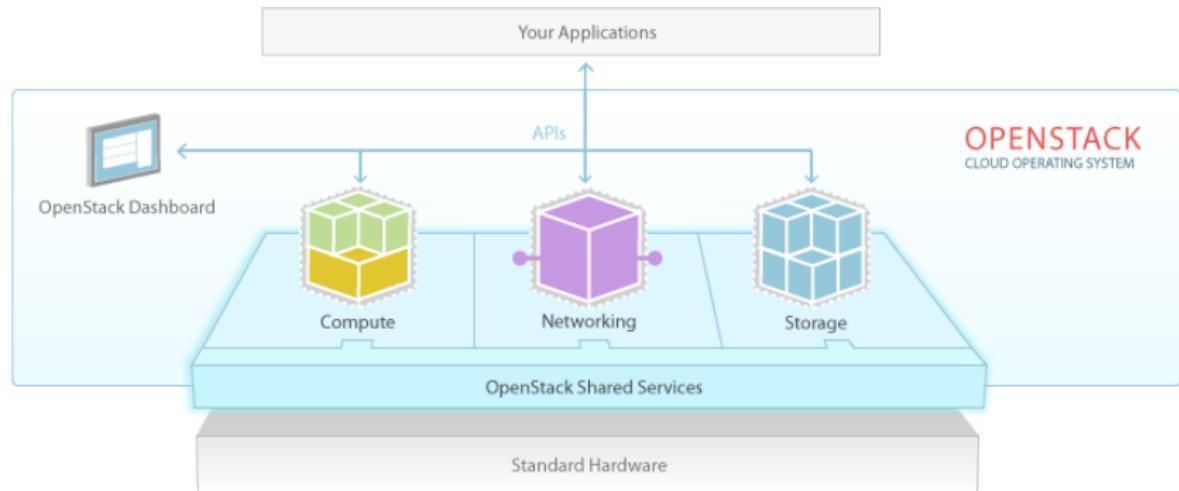
OpenStack is an open source IaaS initiative for creating and managing large groups of virtual private servers in a data center.

The goals of the OpenStack initiative are to

- support interoperability between cloud services
- build cloud services in own data centers



# In a diagram





# History

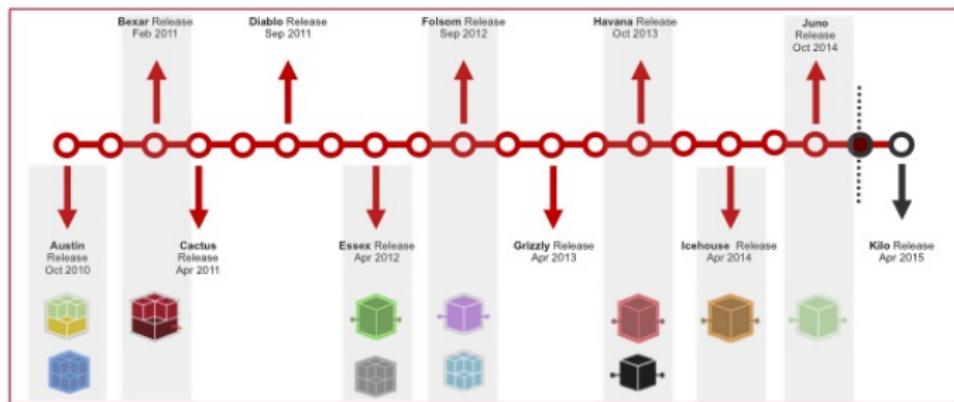
- 2010 : Rackspace and NASA joint effort
- until 2012 : managed by Rackspace and 25 partners
- 2012 : non-profit foundation named OpenStack

Now 600 supporting companies :

<https://www.openstack.org/foundation/companies/>

# OpenStack Releases

## OpenStack® Project Timeline



**rackspace**  
the #1 managed cloud company

[www.rackspace.com](http://www.rackspace.com) 12

and then **Liberty** (2015.2) and **Mitaka** (2016.1)  
next will be **Newton** (2016.2) and **Ocata** (2017.1)



# Services

## Core services :

Nova manages lifecycle of compute instances

Neutron enables network connectivity

Keystone provides authentication

Swift stores and retrieves any data *object*

Cinder provides persistent *block* storage

Glance stores and retrieves VM disk images

And various optional (also less mature) projects :

- Horizon, Heat, Magnum, Sahara, ...



# Community

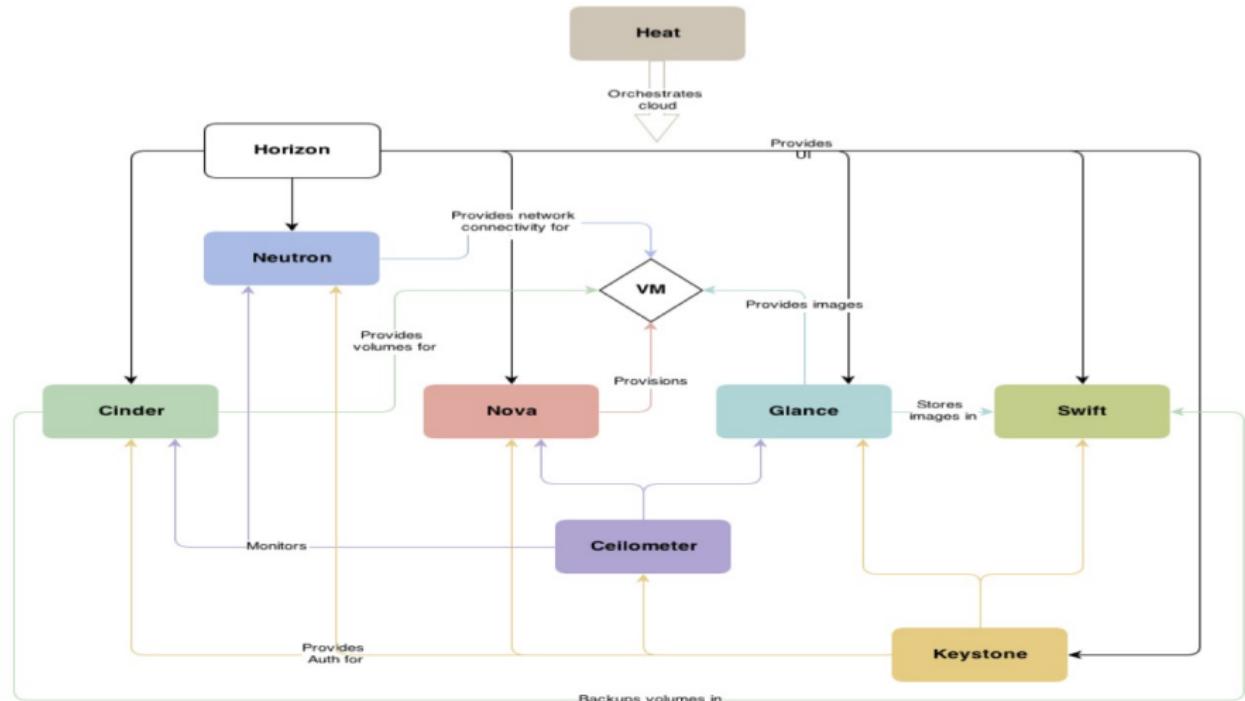
A huge community :

- 40k registered people from 179 countries
- 20M+ lines of code

How it works :

<https://www.openstack.org/software/roadmap/>

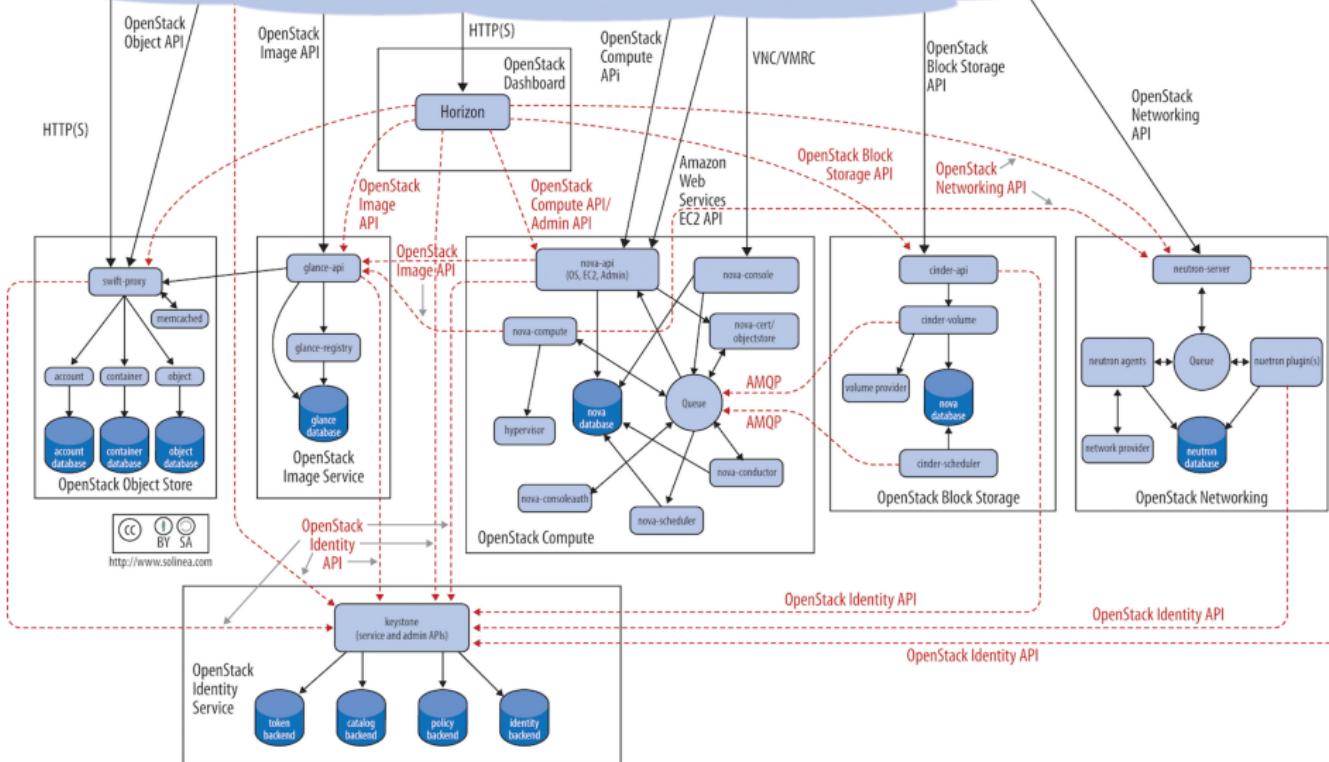
# A more detailed view of services





- Command-line interfaces (nova, neutron, swift, etc)
- Cloud Management Tools (Rightscale, Enstratus, etc)
- GUI tools (Dashboard, Cyberduck, iPhone client, etc)

Internet





# Avancement

- 1 Basic Elements of a Datacenter
- 2 Case Studies
- 3 Dealing with Concerns
- 4 Software Infrastructure
- 5 Conclusion



## A 20-sec conclusion

A new approach of delivering services

A lot of technologies, pushed to their limits